

Cengine

Asynchronous C++ CPU/GPU compute engine, v0.0

Risi Kondor

Flatiron Institute NYC, and
The University of Chicago

Contents

Overview	3
Installation	3
Usage	4
Operators	5
Batched operators	6
Built-in types	7
Reference	8
Built-in types and operators	9
Rscalar	10
Cscalar	11
Ctensor	13
Wrapper classes	15

Overview

Cengine is a lightweight compute engine designed to run in the background, below user level C++ code. The purpose of **Cengine** is to dynamically parallelize computations by

- (a) distributing certain numerical operations across multiple CPU threads and/or
- (b) batching together operations of the same kind for parallel execution on the GPU.

Cengine employs the delayed execution model of computation to decouple user code from the worker threads. It must therefore maintain an internal dependency graph between operations to ensure correctness.

From the user code side, the engine expects a sequence of simple instructions. For example,

```
c=engine.push<ctensor_add_op>(a,b)
```

tells the engine to add tensors **a** and **b** and store the result in **c**. Instead of executing this instruction directly, **Cengine** just adds the corresponding **ctensor_add_op** operator to its internal queue, and executes it later, when either a CPU thread becomes available or a sufficient number of operations of the same type have accumulated to make executing them on the GPU together economical.

Cengine offers a small collection of simple built-in data types such as **rscalar**, **cscalar** and **ctensor** to store real/complex valued scalar and tensor objects, and a corresponding complement of basic arithmetic and linear algebra operators. However, the engine is designed to manage user defined data types and operators just as well as it can manage its built-in types.

Cengine is written in standard C++11 and requires no special libraries besides CUDA/CUBLAS for GPU functionality.

Usage

The most important class in `Cengine` is the compute engine itself, `Cengine::Cengine`. Normally a single instance of this class is initialized at startup and keeps running until the program terminates. For example,

```
Cengine::Cengine engine(4)
```

initializes a compute engine with 4 CPU threads.

User level code does not have direct access to the data objects managed by the engine. Rather, it issues commands to the engine via the engine's `push` method. The `push` method returns a pointer to a handle, which can subsequently be used to reference the object stored internally. For example,

```
Chandle* A=engine.push<new_ctensor_gaussian>({3,3});
```

instructs the engine to create a new 3×3 complex matrix filled with random numbers drawn IID from the standard normal distribution. Issuing the command

```
Chandle* B=engine.push<ctensor_add>(A,A);
```

adds `A` to itself and stores the result in an object referenced by `B`.

Asynchronous execution

`Cengine` follows the asynchronous, delayed execution model of computation. This means that most commands are not executed immediately, but at some later moment in time, when, depending on context, either a CPU thread becomes available, or a sufficient number of operations of the same type have accumulated for execution on the GPU. The order in which the operations are executed need not be the same as the order that they were issued to the engine. To ensure correctness, the engine needs to keep track of dependencies between operations in the form of a directed acyclic graph (DAG).

Delayed execution implies that the `Chandle` objects returned by the engine do not point to the actual result of the computation, but only to where the result will eventually appear. To correctly manage this, user level code will typically encapsulate `Cengine` calls in a separate set of classes. For example, the user might define a `ComplexMatrix` class which has a member variable `hdl` to store the handle returned by engine. To implement in-place matrix addition, the user will add a member function

```
ComplexMatrix& ComplexMatrix::operator+(const ComplexMatrix& B){
    Chandle* t=engine.push<ctensor_add>(hdl,B.hdl);
    delete hdl;
    hdl=t;
}
```

Of course eventually the result of the computation does have to be extracted from the engine. For this we use commands that are *blocking*, meaning that the calling function will wait until the result has actually been computed. For example, the command

```
Gtensor M=engine.push<ctensor_get>(B);
```

returns the value of `B` in a user side tensor object `M`. This command requires explicitly materializing `B`, therefore it waits until all computations leading up to `B` are complete and `B` has been computed as well.

Calling

```
    cengine::flush(B);
```

has a similar effect, while

```
    cengine::flush();
```

flushes all pending operations.

Cengine automatically takes care of memory management. For any given backend object **xobj**, when there are no operations pending that take **xobj** as an argument and no user side handles pointing to **xobj**, the object is scheduled for deletion. When the **Cengine** is deleted or shut down, all pending operations are flushed and all backend objects are destroyed.

Operators

Commands in **Cengine** are implemented as operators, and each command must have a corresponding class. The template argument of **Cengine::push** command is the name of this operator class. The abstract base class of all operator classes is **Cengine::Coperator**. For example, the internal definition of the **ctensor_add_op** operator (slightly abbreviated) is

```
class ctensor_add_op: public Coperator, public CumulativeOperator{
public:

    using Coperator::Coperator;

    void exec(){
        owner->obj=inputs[0]->obj;
        asCtensorB(owner).add(asCtensorB(inputs[1]));
    }

    string str() const{
        return "ctensor_add"+inp_str();
    }

};
```

In each operator class, the **exec** method is responsible for carrying out the actual operation on the operator's arguments. In the case of **ctensor_add_op** this just amounts to calling the backend object's method for tensor summation. However, in user-defined operators the **exec** method can often be significantly more involved. Adding a new operator to **Cengine** just requires defining the corresponding operator class.

The concrete data object that **ctensor_add_op** operates on is a **Cengine::CtensorB**, which is the backend container for **ctensor** objects. The built-in **rscalar**, **scalar** and **ctensor** classes, extended by user defined operators are sufficient for many purposes. However, there is nothing stopping the user from adding new backend classes as well, simply by subclassing the **Cengine::Cobject** abstract class. **Cengine** can manage any type of user defined backend object, as long as it is derived from **Cengine::Cobject**, and any type of user defined operator, as long as it is derived from **Cengine::Coperator**. Adding new objects and operators does not require making any changes to **Cengine** itself.

Batched operators

Batching refers to accumulating multiple instances of the same operation and executing them together, in parallel. Batching is particularly important efficiently utilizing graphics processor units (GPUs), since GPU threads are generally tied: on NVIDIA architectures, for example, all threads running on the same streaming multiprocessor must essentially be executing the the exact same machine level instruction at any given time. Some types of computations, such as solving a systems of partial differential equations on a regular grid are well suited to this paradigm, since the operations that need to be performed at each gridpoint are the same.

Other types of computations, however, are much less structured. In a graph neural network, for example, the operation performed at each node depends on the number of neighbors. In principle, it is possible to write code that separately parallelizes over all nodes with just one neighbor, all nodes with two neighbors, and so on, but in practice such low level multithreading is laborious and highly error prone.

One solution that has emerged is *dynamic batching*, which refers to accumulating operations of each type and executing them together as a batch. Taking dynamic batching too far can lead to situations where a large number of batched operations are mutually waiting on each other and none of the batches are actually run. Therefore, as a general principle, it is best to use dynamic batching sparingly, on a small set of frequent operations that are expensive enough to be performance critical, yet small enough that executing the operations individually (without batching) would waste much of the GPU's parallel processing power. Basic matrix operations, such as matrix/scalar and matrix/matrix products are good candidates for batching. Accessing individual components of matrices, however, is not an operation that would likely benefit from dynamic batching.

Cengine will attempt to dynamically batch any operator derived from the **BatchedOperator** class. In addition to the **exec()** method, batched operators must also have a **batched.exec** method, which takes a *vector* of pointers to compute graph nodes as its argument, and executes each node in parallel. For each batched operator class **UserOp1**, the engine will internally create a separate **BatcherA<UserOp1>** object to manage the batching process. To keep track of the correspondence between operators and batchers, each batched operator class must provide a static integer **batcher_id** variable.

Meta-batchers

Many batched operators require separate batchers for different settings of their parameters. For example, in order to batch matrix multiplication, we need separate batchers for each combination of input matrix dimensions. **Cengine** makes it easy to implement such *multi-batched* operators by introducing batcher signatures and the **MetaBatcher** class.

Any multi-batched class must have a corresponding signature type. For example the signature class of the matrix multiplication operator is **Mprod.signature**, which stores the dimensions of the two matrices to be multiplied and some flags to signify if either matrix is transposed. The matrix multiplication operator **ctensor.Mprod_op** must have a **signature()** method that returns the signature object corresponding to the given pair of matrices to be multiplied. The engine will then create a separate batcher object for each distinct signature.

The purpose of the **MetaBatcher** class (for matrix multiplication) is to route individual matrix products to their corresponding batcher. All that the operator class needs to do enable this process is provide a **spawn_batcher()** method that creates the appropriate templated **MetaBatcher** object. In the case of our example the type of this (in slightly abbreviated form) would be

```
MetaBatcher< ctensor_add_Mprod_op, Mprod.signature, BatcherA<ctensor_add_Mprod_op> >.
```

Built-in types

While **Cengine** is primarily designed to be used with user-defined data classes and operators, it has a small collection of simple built-in types corresponding to real and complex scalars/tensors to seed this process:

```
rscalar  
cscalar  
ctensor .
```

To enable fast GPU computation, each of these classes is implemented in single precision arithmetic (**float**). The corresponding back-end classes are **RscalarB**, **CscalarB** and **CtensorB**. Each of these types is equipped with a minimal set of arithmetic and linear algebra operators.

Data layout and GPU functionality

Similarly to deep learning frameworks such as **PyTorch** and **TensorFlow**, **Cengine**'s built in objects can be flexibly moved back and forth between the host and the GPUs. This is done by the **to_device(d)** command, where **d** is the identifier of the GPU, or **0** in case that the object is to be moved back to the host.

In general, every backend operation has two separate implementations: one for the execution on the CPU in straight-line C++, and once for execution on the GPU written in CUDA or CUBLAS. Whether a given operation is executed on the CPU or GPU depends on where its arguments reside: in general, **Cengine** will move all input objects to the same device as where the first input argument resides and perform the operation on that device.

The storage layout of the built in classes is optimized for GPU computation. In particular, matrix/tensor objects are padded to multiples of 32 **floats**, and complex tensor are stored with their real and imaginary parts separate. **Cengine** uses a row-major matrix/tensor storage format.

Reference

Built-in types and operators

Rscalar

CONSTRUCTORS

`new_rscalar_op(const int nbu=-1, const int dev=0)`

`new_rscalar_zero_op(const int nbu=-1, const int dev=0)`

`new_rscalar_gaussian_op(const int nbu=-1, const int dev=0)`

Construct a new `rscalar` object with `nbu` bundles on device `dev`. The three cases correspond to the object being uninitialized, initialized to zero, or initialized with random standard normal entries. `nbu=-1` signifies that the object is not bundled and `dev=0` is the host.

`rscalar_copy_op(const rscalar& x)`

Create a new `rscalar` by copying `x`.

IN-PLACE OPERATORS

`rscalar_set_zero_op(const rscalar& x)`

Set `x` to zero.

CUMULATIVE OPERATORS

`rscalar_add_op(const rscalar& r, const rscalar& x)`

Set $r \leftarrow r + x$.

`rscalar_subtract_op(const rscalar& r, const rscalar& x)`

Set $r \leftarrow r - x$.

`rscalar_add_prod_op(const rscalar& r, const rscalar& x, const rscalar& y)`

Set $r \leftarrow r + xy$.

`rscalar_add_div_op(const rscalar& r, const rscalar& x, const rscalar& y)`

Set $r \leftarrow r + x/y$.

`rscalar_add_abs_op(const rscalar& r, const rscalar& x)`

Set $r \leftarrow r + |x|$.

`rscalar_add_exp_op(const rscalar& r, const rscalar& x)`

Set $r \leftarrow r + e^x$.

`rscalar_add_pow_op(const rscalar& r, const rscalar& x, const float p, const float c)`

Set $r \leftarrow r + cx^p$.

BLOCKING OPERATORS

Cscalar

CONSTRUCTORS

`new_cscalar_op(const int nbu=-1, const int dev=0)`

`new_cscalar_zero_op(const int nbu=-1, const int dev=0)`

`new_cscalar_gaussian_op(const int nbu=-1, const int dev=0)`

Construct a new `cscalar` object with `nbu` bundles on `dev`. The three cases correspond to the object being uninitialized, initialized to zero, or initialized with random standard normal entries. `nbu=-1` signifies that the object is not bundled and `device=0` is the host.

`cscalar_copy_op(const cscalar& x)`

Create a new `cscalar` by copying `x`.

IN-PLACE OPERATORS

`cscalar_set_zero_op(const cscalar& x)`

Set `x` to zero.

NOT IN-PLACE OPERATORS

`cscalar_conj_op(const cscalar& x)`

Return \bar{x} .

`cscalar_get_real_op(const cscalar& x)`

Return the real part of x .

`cscalar_get_imag_op(const cscalar& x)`

Return the imaginary part of x .

CUMULATIVE OPERATORS

`cscalar_add_op(const cscalar& r, const cscalar& x)`

Set $r \leftarrow r + x$.

`cscalar_subtract_op(const cscalar& r, const cscalar& x)`

Set $r \leftarrow r - x$.

`cscalar_add_prod_op(const cscalar& r, const cscalar& x, const cscalar& y)`

Set $r \leftarrow r + xy$.

`cscalar_add_div_op(const cscalar& r, const cscalar& x, const cscalar& y)`

Set $r \leftarrow r + x/y$.

`cscalar_add_to_real_op(const cscalar& r, const rscalar& x)`

Set $r \leftarrow r + (x, 0)$.

`cscalar_add_to_imag_op(const cscalar& r, const rscalar& x)`

Set $r \leftarrow r + (0, x)$.

`cscalar_add_abs_op(const cscalar& r, const cscalar& x)`

Set $r \leftarrow r + |x|$.

```
cscalar_add_exp_op(const cscalar& r, const cscalar& x)
```

Set $r \leftarrow r + e^x$.

```
cscalar_add_pow_op(const cscalar& r, const cscalar& x, const float p, const float c)
```

Set $r \leftarrow r + cx^p$.

Ctensor

CONSTRUCTORS

`new_ctensor_op(const Gdims& dims, const int nbu=-1, const int dev=0)`

`new_ctensor_zero_op(const Gdims& dims, const int nbu=-1, const int dev=0)`

`new_ctensor_identity_op(const Gdims& dims, const int nbu=-1, const int dev=0)`

`new_ctensor_gaussian_op(const Gdims& dims, const int nbu=-1, const int dev=0)`

Construct a new `ctensor` object of size `dims` with `nbu` bundles on `dev`. The four cases correspond to the object being uninitialized, initialized to zero, the identity matrix, or initialized with random standard normal entries. `nbu=-1` signifies that the object is not bundled and `device=0` is the host.

`new_ctensor_from_gtensor_op(const Gtensor& T, const int nbu=-1, const int dev=0)`

Create a new `ctensor` from the `Gtensor` `T`.

`ctensor_copy_op(const ctensor& x)`

Create a new `ctensor` by copying `x`.

IN-PLACE OPERATORS

`ctensor_set_zero_op(const ctensor& x)`

Set `x` to zero.

NOT IN-PLACE OPERATORS

`ctensor_conj_op(const ctensor& x)`

Return \bar{x} .

`ctensor_get_real_op(const ctensor& x)`

Return the real part of x .

`ctensor_get_imag_op(const ctensor& x)`

Return the imaginary part of x .

CUMULATIVE OPERATORS

`ctensor_add_op(const ctensor& r, const ctensor& x)`

Set $r \leftarrow r + x$.

`ctensor_add_conj_op(const ctensor& r, const ctensor& x)`

Set $r \leftarrow r + \bar{x}$.

`ctensor_add_transp_op(const ctensor& r, const ctensor& x)`

Set $r \leftarrow r + x^{\top}$.

`ctensor_add_herm_op(const ctensor& r, const ctensor& x)`

Set $r \leftarrow r + x^{\dagger}$.

`ctensor_add_times_real_op(const ctensor& R, const ctensor& A, const float c)`

Set $R \leftarrow R + cA$.

`ctensor_add_times_complex_op(const ctensor& R, const ctensor& A, const complex<float> c)`

Set $R \leftarrow R + cA$.

ctensor_add_prod_rA.op(const ctensor& R, const rscalar& c, const ctensor& A)

Set $R \leftarrow R + cA$.

ctensor_add_prod_cA.op(const ctensor& R, const cscalar& c, const ctensor& A)

Set $R \leftarrow R + cA$.

ctensor_add_Mprod.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_AT.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_TA.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_AC.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_TC.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_AH.op(const ctensor& R, const ctensor& A, const ctensor& B)

ctensor_add_Mprod_HA.op(const ctensor& R, const ctensor& A, const ctensor& B)

Set $R \leftarrow R + AB$, $R \leftarrow R + AB^\top$, $R \leftarrow R + A$, $R \leftarrow R + A\bar{B}$, $R \leftarrow R + A^\top \bar{B}$, $R \leftarrow R + AB^\dagger$, $R \leftarrow R + A$.

INTO OPERATORS

ctensor_add_inp.op(const cscalar& r, const ctensor& A, const ctensor& B)

Set $r \leftarrow r + \langle A, B \rangle$.

Wrapper classes